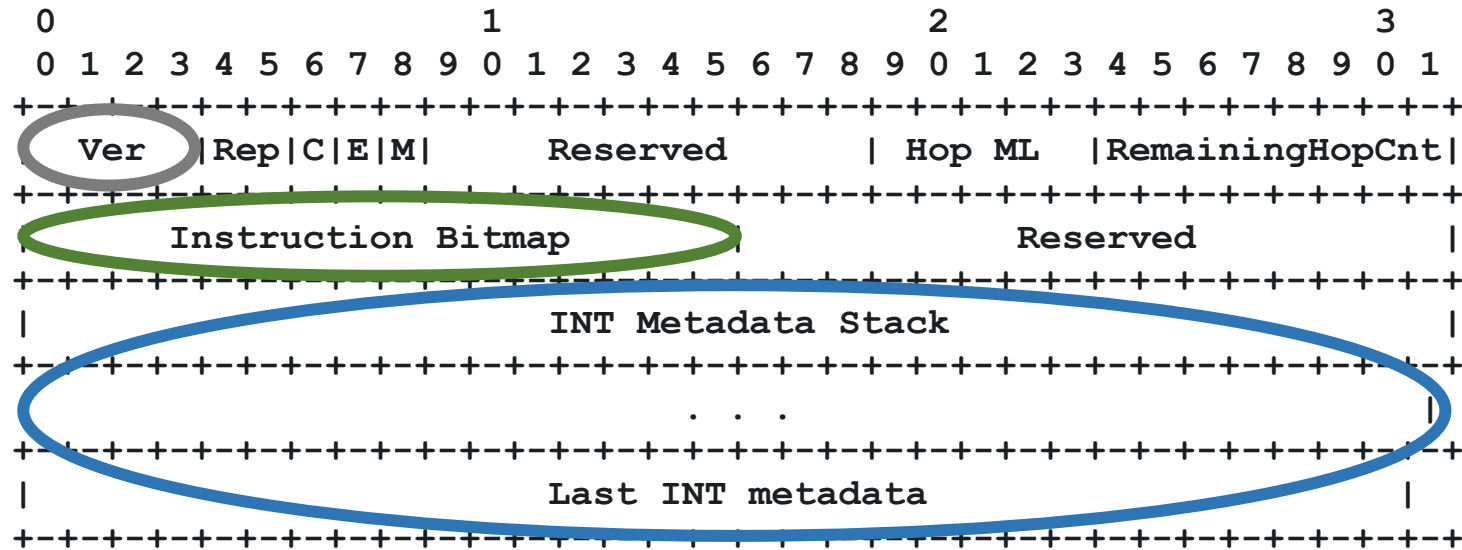


Conveying Metadata Semantics to Monitors

P4 Apps working group

May 17, 2018

INT v1.0 Hop-by-Hop Metadata Header Format



INT instructions are encoded as a bitmap in the 16-bit INT Instruction field:

bit0 (MSB): Switch ID

bit1: Level 1 Ingress Port ID (16 bits) + Egress Port ID (16 bits)

bit2: Hop latency

bit3: Queue ID (8 bits) + Queue occupancy (24 bits)

bit4: Ingress timestamp

bit5: Egress timestamp

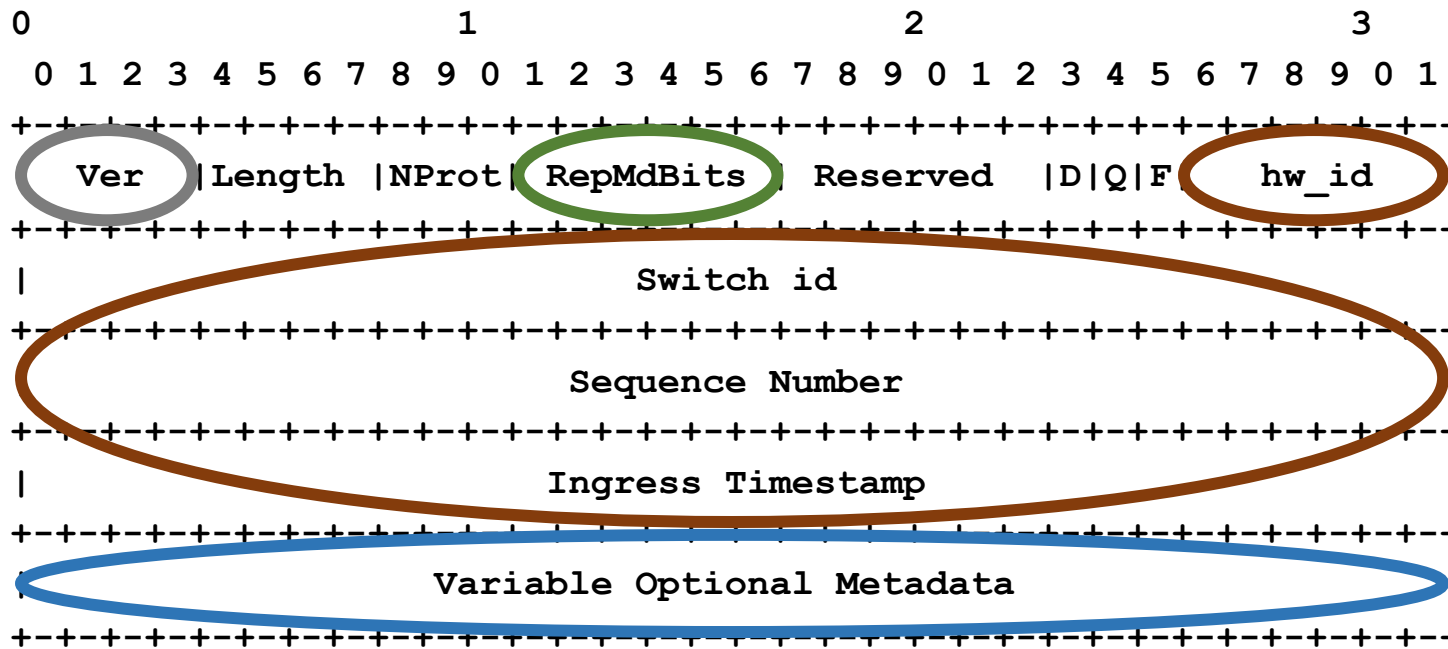
bit6: Level 2 Ingress Port ID + Egress Port ID (4 bytes each)

bit7: Egress port Tx utilization

bit15: Checksum Complement

The remaining bits are reserved.

Telemetry Report v1.0 Header



RepBits: Report Metadata Bits indicate which optional metadata (4 octets each) is present

bit 0 (MSB): Ingress port id (16 bits) + Egress port id (16 bits)

bit 1: Hop latency

bit 2: Queue id (8 bits) + Queue occupancy (24 bits)

bit 3: Egress Timestamp (32 bits)

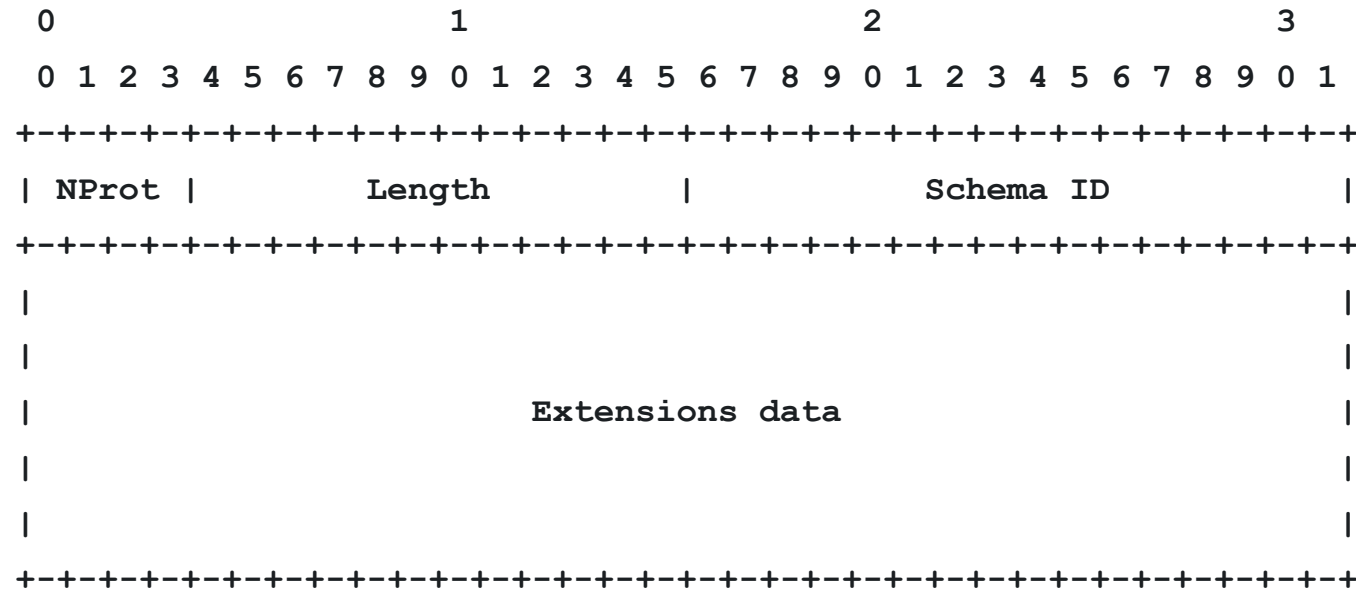
bit 4: Queue id (8 bits) + Drop reason (8 bits) + Padding (16 bits)

bit 5: Egress port tx utilization

Aspects of Metadata Semantics

- Supported or not supported
 - May vary by version
 - May vary for INT metadata vs Telemetry Report metadata
- Units
 - Queue occupancy in bytes, cells, or packets?
Cell size?
 - Hop latency in nanoseconds, microseconds, or $2^{(-32)}$ seconds?
 - Utilization in percent?
 - Drop reason enum definition? By reference?
- Boundary locations
 - Queue occupancy at enqueue time or at dequeue time? Includes this packet?
 - Hop latency is from where to where?
- Anything to say about ID definitions?
 - Switch
 - Port levels 1 + 2
 - Queue
- Report location for last hop metadata
 - embedded INT
 - telemetry report metadata in same report, or
 - telemetry report metadata in separate report

Proposed Telemetry Extensions Header



Presence is indicated by new *NProt* codepoint in Telemetry Report Header

- NProt: Next Protocol: Same definition as *NProt* in the Telemetry Report Header.
- Length: Indicates the length of the telemetry extensions header in multiples of 4 octets.
- Schema ID: 2-octet unsigned integer identifying the schema of the telemetry extensions data:
 - 0 - 0xEFFF : Available for private or experimental use.
 - 0xF000 - 0xFEFF : Reserved.
 - 0xFF00 - 0xFFFE : Reserved for specification by P4.org. This range is set aside for future uses such as specification of field units in INT and telemetry reports.
 - 0xFFFF : No schema, indicates that the telemetry extensions header has no content.
- Extensions data: Variable length field. This field is interpreted as specified by the schema identified by the Schema ID.

Encoding Within Telemetry Extensions

- Generated by CPU on switch, consumed by monitor
→ ease of implementation in hardware is not a concern
- Encoding options
 - List of fixed-length key value pairs
 - TLVs
 - Strings, e.g. JSON
- Dedicate one entire Schema ID value for metadata semantics?

Encoding using fixed-length key value pairs

Key	Meaning	Value Syntax	Comments
1	Hop latency	enum { nanoseconds, microseconds, 2 ⁽⁻³²⁾ seconds }	
2	Queue occupancy	enum { bytes, cells, packets }	If cells, what is the cell size? Subdivide value? Or cell size is part of enum definition? e.g. 80 byte cells, 100 byte cells, ...
3	Tx utilization	enum { percent }	

- 16 bit key + 16 bit value for each key value pair?
- How to convey semantics other than units?
Separate key value pair for queue occupancy location?
- Is there anything that requires a value syntax other than enum?

Encoding using nested TLVs

Type			Meaning	Length	Value Syntax
1			Metadata	6	pad
	1		Hop latency	1	enum { nanoseconds, microseconds, 2 ⁽⁻³²⁾ seconds }
	2		Queue occupancy	3	enum { bytes, cells, packets }
		1	Cell size	1	integer
		2	Packet timing	1	enum { enqueue, dequeue }
	3		Tx utilization	1	enum { percent }
2			Last hop metadata	1	enum { embedded INT, same telemetry report, different telemetry report }

1 byte type

1 byte length in units of 4 octets
(inclusive)

2 + 4*n bytes fixed part of value
arbitrary nested TLVs

Encoding using JSON (JSON Schema)

```
{
  "metadata" : {
    "type" : "object",
    "properties" : {
      "hop latency" : {
        "type" : "object",
        "properties" : {
          "units" : { "type" : "string",
                     "enum" : ["nanoseconds", "milliseconds", "2(-32) seconds"] }
        }
      },
      "queue occupancy" : {
        "type" : "object",
        "properties" : {
          "units" : { "type" : "string",
                     "enum" : ["bytes", "cells", "packets"] },
          "cell size" : { "type" : "integer" },
          "packet timing" : { "type" : "string",
                              "enum" : ["enqueue", "dequeue"] }
        }
      }
    }
  },
  "last hop metadata" : { "type" : "string",
                          "enum" : ["embedded INT", "same telemetry report", "different telemetry report"] }
}
```

Encoding using JSON (JSON example)

```
{
  "metadata" : {
    "hop latency" : {
      "units" : "nanoseconds"
    },
    "queue occupancy" : {
      "units" : "cells"
      "cell size" : 80
      "packet timing" : "enqueue"
    },
    "hop latency" : {
      "units" : "percent"
    }
  },
  "last hop metadata" : "same telemetry report"
}
```